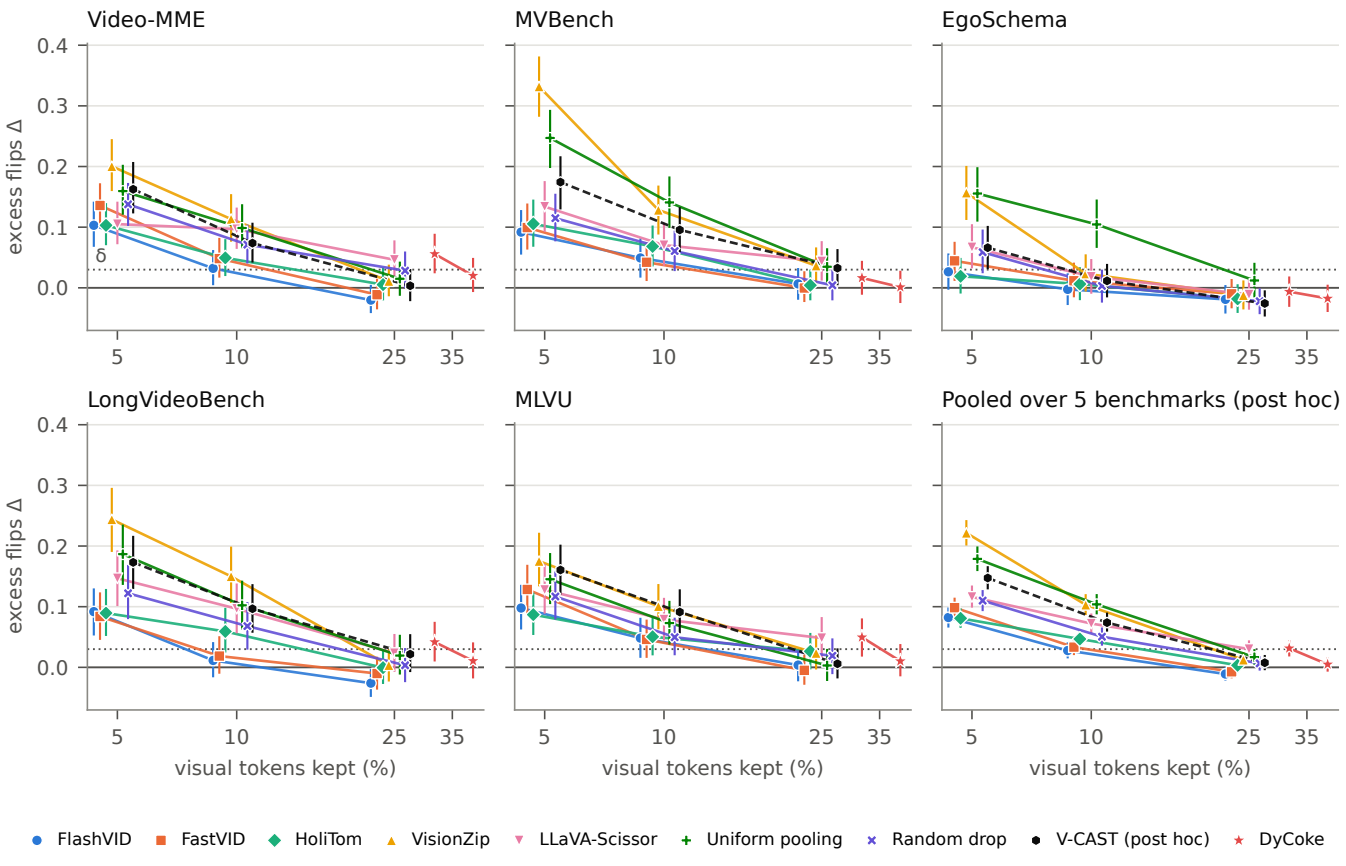
Video-MME
MVBench
EgoSchema
LongVideoBench
MLVU
Pooled over 5 benchmarks (post hoc)
excess flips Δ
δ
0.4
0.3
0.2
0.1
0.0
5
10
25
35
visual tokens kept (%)
FlashVID
FastVID
HoliTom
VisionZip
LLaVA-Scissor
Uniform pooling
Random drop
V-CAST (post hoc)
DyCoke